

[]: Summary :

Training Accuracy : 0.99

Testing Accuracy: 0.99

No. of iteration to converge : 200

Do you think your model overfits?

-> No the model does **not** over fit

Features **with** highest absolute weight :

-> 'dst_host_srv_error_rate', 'error_rate', 'dst_host_rerror_rate'

Show how the train **and** test accuracy varies **with** different regularization
strengths.

-> Different regularizations **and** solver combination tried using 'penalty = '
and 'solver = '.

With the increase **in** strength , the convergence decreases

```
[14]: import numpy as np
import pandas as pd
```

```
features =
```

```
↳ ['duration', 'protocol_type', 'service', 'flag', 'src_bytes', 'dst_bytes', 'land', 'wrong_fragment', 'u
'hot', 'num_failed_logins', 'logged_in', 'num_compromised', 'root_shell', 'su_attempted', 'num_root', 'n
'num_shells', 'num_access_files', 'num_outbound_cmds', 'is_host_login', 'is_guest_login', 'count', 'srv
'serror_rate', 'srv_error_rate', 'rerror_rate', 'srv_rerror_rate', 'same_srv_rate', 'diff_srv_rate', 's
'dst_host_count', 'dst_host_srv_count', 'dst_host_same_srv_rate', 'dst_host_diff_srv_rate', 'dst_host
'dst_host_srv_diff_host_rate', 'dst_host_serror_rate', 'dst_host_srv_error_rate', 'dst_host_rerror_r
'label']
```

```
#intrusion_type = ['back', 'buffer_overflow', 'ftp_write',
# 'guess_passwd', 'imap', 'ipsweep', 'land', 'loadmodule',
# 'multihop', 'neptune', 'nmap', 'normal', 'perl',
# 'phf', 'pod', 'portsweep', 'rootkit',
# 'satan', 'smurf', 'spy', 'teardrop',
# 'warezclient', 'warezmaster']
```

```
[15]: data = pd.read_csv('kddcup.data', names = features, header=None)
data.drop(data.loc[data['label']=='normal.'].index, inplace=True)
```

Converting the label probe/non probe to 0 and 1

```
[16]: l_nprobe = ['back.', 'buffer_overflow.', 'ftp_write.', 'guess_passwd.', 'imap.', 'land.
↳ ', 'loadmodule.', 'multihop.', 'neptune.', 'perl.', 'phf.', 'pod.', 'rootkit.', 'smurf.
↳ ', 'spy.', 'teardrop.', 'warezclient.', 'warezmaster.
l_probe = ['ipsweep.', 'nmap.', 'portsweep.', 'satan.']
```

```
#0 if normal , 1 if non-probe , 2 if probe

data['label'] = data['label'].apply(lambda x: 2 if x == 'normal.' else (0 if (x in
↳ l_nprobe ) else 1))

# Get unique labels : print(data['label'].unique())
# Review the columns which might need integer encoding data.dtypes
```

```
[17]: from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score

from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import make_column_transformer
```

One Hot Encoding the columns

```
[18]: protocol_end = pd.get_dummies(data.protocol_type)
service_end = pd.get_dummies(data.service)
flag = pd.get_dummies(data.flag)
```

```
[27]: data_encoded = pd.concat([data,protocol_end,service_end,flag],axis = 'columns')
```

```
[29]: data_encoded_final = data_encoded.drop(['protocol_type','flag','service'],axis =
↳ 'columns')
#data_encoded.drop('service',axis = 'columns')
#data_encoded.drop('flag',axis = 'columns')
```

```
le = LabelEncoder() protocol_end = le.fit_transform(data['protocol_type']) data.drop('protocol_type',axis
= 'columns') data['protocol_type'] = protocol_end
```

```
le1 = LabelEncoder() service_end = le1.fit_transform(data['service']) data.drop('service',axis = 'columns')
data['service'] = service_end
```

```
le2 = LabelEncoder() flag_end = le2.fit_transform(data['flag']) data.drop('flag',axis = 'columns')
data['flag'] = flag_end
```

if validation of what all categories were mapped is needed

```
#le1.classes_ #le2.classes_
```

```
#Quick look at how the data looks #data.head(10) #data.dtypes #data.drop_duplicates(subset=features,
keep='first', inplace = True)
```

```
#Save the data to a pickle file to avoid reprocessing and enable restartability #data.to_pickle('data.pkl')
```

X and y being created

```
[31]: X = data_encoded_final.drop('label',axis = 'columns')
y = data_encoded_final.label

# validate if X and y have same rows # y.shape
```

```
[33]: y.shape
```

```
[33]: (3925650,)
```

Test and Training Split and Scaling the data

```
[34]: from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import StandardScaler

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=1, stratify=y)

sc = StandardScaler()
sc.fit(X_train)

X_train_std = sc.transform(X_train)
X_test_std = sc.transform(X_test)

print('Labels count in y:', np.bincount(y))
print('Labels count in y_train:', np.bincount(y_train))

print('Labels count in y_test:', np.bincount(y_test))
```

```
Labels count in y: [3884548  41102]
```

```
Labels count in y_train: [3107638  32882]
```

```
Labels count in y_test: [776910   8220]
```

Logistic Regression using Scikit learn

```
[40]: logreg = LogisticRegression(C=1e10, solver='lbfgs', multi_class='ovr', penalty = 'l2',
    ↪, max_iter = 1000, random_state = 1)
logreg.fit(X_train_std, y_train)

#logreg.fit(X_train_std, y_train)
```

```
/Users/samipsinghal/opt/anaconda3/lib/python3.7/site-
packages/sklearn/linear_model/logistic.py:947: ConvergenceWarning: lbfgs failed to
converge. Increase the number of iterations.
"of iterations.", ConvergenceWarning)
```

```
[40]: LogisticRegression(C=10000000000.0, class_weight=None, dual=False,
    fit_intercept=True, intercept_scaling=1, l1_ratio=None,
    max_iter=1000, multi_class='ovr', n_jobs=None, penalty='l2',
    random_state=1, solver='lbfgs', tol=0.0001, verbose=0,
    warm_start=False)
```

Training and Test Error

```
[37]: y_pred = logreg.predict(X_test_std)
      print('Misclassified examples: %d' % (y_test != y_pred).sum())
```

Misclassified examples: 12

```
[38]: from sklearn.metrics import accuracy_score
      print('Accuracy: %.3f' % accuracy_score(y_test, y_pred))
```

Accuracy: 1.000

Number of Iterations to Converge

```
[39]: print(logreg.n_iter_)
```

[200]

Plotting the features with highest absolute weight

```
[41]: import numpy as np
      import matplotlib.pyplot as plt

      # We create a matrix with all the labels
      x_labels_w = X.columns[1:]
      Xw = np.array(X[x_labels_w].values)
      print(Xw)
      print("The matrix dimensions of Xw is " + str(Xw.shape))

[[1.511e+03 2.957e+03 0.000e+00 ... 0.000e+00 1.000e+00 0.000e+00]
 [1.735e+03 2.766e+03 0.000e+00 ... 0.000e+00 1.000e+00 0.000e+00]
 [2.810e+02 1.301e+03 0.000e+00 ... 0.000e+00 1.000e+00 0.000e+00]
 ...
 [2.800e+01 0.000e+00 0.000e+00 ... 0.000e+00 1.000e+00 0.000e+00]
 [2.800e+01 0.000e+00 0.000e+00 ... 0.000e+00 1.000e+00 0.000e+00]
 [2.800e+01 0.000e+00 0.000e+00 ... 0.000e+00 1.000e+00 0.000e+00]]
The matrix dimensions of Xw is (3925650, 117)
```

```
[42]: # For plotting in the Jupyter Notebook environment as an inline output
      %matplotlib inline
```

```
[43]: # By default, LogisticRegression() is set on penalty as L2 and C=1.
      # To simulate no regularization, we will select a large C to minimize regularization
      ↪to later
      # show the effect of regularization

      #logreg_w=LogisticRegression(C=1e10,solver = 'lbfgs')

      #logreg_w.fit(Xw,y)


      W=logreg.coef_

      W=W.flatten()
```

```
plt.stem(W)
```

/Users/samipsinghal/opt/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:12:
UserWarning: In Matplotlib 3.3 individual lines on a stem plot will be added as a
LineCollection instead of individual lines. This significantly improves the performance
of a stem plot. To remove this warning and switch to the new behaviour, set the
"use_line_collection" keyword argument to True.
if sys.path[0] == '':

```
[43]: <StemContainer object of 3 artists>
```



```
[44]: idx1=np.argsort(np.abs(W))[-1]
      idx2=np.argsort(np.abs(W))[-2]
      idx3=np.argsort(np.abs(W))[-3]

      heavy=[x_labels_w[idx1], x_labels_w[idx2],x_labels_w[idx3]]
      heavy
```

```
[44]: ['dst_host_srv_serror_rate', 'serror_rate', 'dst_host_rerror_rate']
```

```
[ ]:
```